

Devices and circuits for Artificial Intelligence

-- 540438 - Maurizio La Rosa --

Final exam project

In this notebook I introduce the project for the final exam of the course *Devices and circuits for artificial intelligence* from the Data Analysis degree of the University of Messina. The project consists in building a machine learning model for image classification.

The dataset to be used is hosted at [kaggle](#), and is called [BIRDS 525 SPECIES- IMAGE CLASSIFICATION](#). The dataset currently contains images from 525 bird species to be classified by the model. I downloaded the dataset on April, 17th, 2023, and that version contains 515 bird species.

It is useful to note that images in the dataset should have all the exact same shape (224, 224, 3), while I found that all images of the 'PLUSH CRESTED JAY' species and one image from the 'DON'T REMEMBER THE SPECIES, FILL INFO WITH FUTURE COMMIT' species have variable shapes. Hence, in my code, I check for images' shapes and remove images that don't match the common shape. This is important because imported images have the shape of 3D Numpy (np, when imported) arrays and I need to transform the list of images into a 4D Numpy array. The function `np.array()` can do it automatically when fed a list of 3D Numpy arrays, but images must have all the same shape.

The following cell is for importing necessary modules in the file used for describing the data, the model and the results. I previously uploaded data for bird species classification on Colab, in the *content/kaggle_data* folder.

```

In [ ]: #####
      ### allow importing from Google Drive ###
      ### after uploading the needed files ###
      #####
      from google.colab import drive
      drive.mount('/content/drive', force_remount=True)
      import sys
      sys.path.insert(0, '/content/drive/MyDrive/Colab Notebooks/da_dcAI_project

      #####
      ### unzip reshaped dataset from Google ###
      ### Drive to colab's content folder ###
      #####
      !unzip '/content/drive/MyDrive/Colab Notebooks/da_dcAI_project/reshaped_d

      #####
      ### import necessary or useful modules ###
      #####
      import os
      import numpy as np
      #import tensorflow as tf
      from matplotlib import pyplot as plt
      from tensorflow.keras.utils import to_categorical
      from keras.preprocessing.image import ImageDataGenerator
      from sklearn.metrics import classification_report, confusion
      from a_selectRandomFolders import selectRandomFolders
      from b_viewClasses import viewRandomClasses, viewRandomCla
      from c_selectData import selectData, classesToInt, countL
      from d_tl_model_ResNet import seqModel
      from e_plotConfusionMatrix import plotCM

```

Retrieve data folders

The following code cell is used for retrieving the system folders where the data are located.

```
In [4]: ### 1) set path to the directory of the dataset (this is the targetFolder
### and show how many bird classes there are in the path and their nam
trainPathWin = 'C:/Users/mzlarosa/OneDrive - unime.it/Learning/CdL Informa
testPathWin = 'C:/Users/mzlarosa/OneDrive - unime.it/Learning/CdL Informa
validPathWin = 'C:/Users/mzlarosa/OneDrive - unime.it/Learning/CdL Informa
trainPathMac = '/Users/mau/OneDrive - unime.it/Learning/CdL Informatica/A
testPathMac = '/Users/mau/OneDrive - unime.it/Learning/CdL Informatica/An
validPathMac = '/Users/mau/OneDrive - unime.it/Learning/CdL Informatica/A
trainPath = '/content/reshaped/train'
testPath = '/content/reshaped/test'
validPath = '/content/reshaped/valid'
classes = os.listdir(trainPath)
if '.DS_Store' in classes:
    classes.remove('.DS_Store')
nOfClasses = len(classes)
print('\nThere are', nOfClasses, 'classes in the dataset.\n' +
      '\nHere is a list of the first 50 classes:')
print(classes[0 : 50], end = '')
print(' [...]')
```

There are 515 classes in the dataset.

Here is a list of the first 50 classes:

```
['BLACK-CAPPED CHICKADEE', 'BLONDE CRESTED WOODPECKER', 'WHITE BROWED CR
AKE', 'AMERICAN GOLDFINCH', 'LIMPKIN', 'APAPANE', 'BLACKBURNIAM WARBLE
R', 'NORTHERN FLICKER', 'PAINTED BUNTING', 'AUSTRALASIAN FIGBIRD', 'ANDE
AN SISKIN', 'EMPEROR PENGUIN', 'TRICOLORED BLACKBIRD', 'RUBY THROATED HU
MMINGBIRD', 'PHILIPPINE EAGLE', 'MALACHITE KINGFISHER', 'GYRFALCON', 'MI
LITARY MACAW', 'GOLD WING WARBLER', 'VARIED THRUSH', 'GROVED BILLED AN
I', 'EASTERN BLUEBONNET', 'MALLARD DUCK', 'NORTHERN MOCKINGBIRD', 'HOUSE
SPARROW', 'OYSTER CATCHER', 'CAPPED HERON', 'ASIAN CRESTED IBIS', 'COMMO
N POORWILL', 'OCELLATED TURKEY', 'IVORY BILLED ARACARI', 'NORTHERN RED B
ISHOP', 'QUETZAL', 'TROPICAL KINGBIRD', 'CRIMSON SUNBIRD', 'AFRICAN CROW
NED CRANE', 'GREAT TINAMOU', 'GREAT KISKADEE', 'FRILL BACK PIGEON', 'GIL
A WOODPECKER', 'HYACINTH MACAW', 'HARLEQUIN DUCK', 'GREY PLOVER', 'LILAC
ROLLER', 'FASCIATED WREN', 'BROWN CREPPER', 'SATYR TRAGOPAN', 'ABBOTTS B
ABBLER', 'CEDAR WAXWING', 'BEARDED BELLBIRD'] [...]
```

Introduce the data

The number of available pictures varies with the class.

I select a random sample of 15 bird species from the train data and show how many images are available for each sampled species.

```
In [ ]: ### 2) select a random sample of n (15) subfolders from the targetFolder
### and show their content (subfolders represent bird classes)
### modules: os, random, selectRandomFolders
targetClasses = selectRandomFolders(trainPath, 15)
```

There are 0 folders and 157 image files in /content/reshaped/train/EGYPTIAN GOOSE
There are 0 folders and 152 image files in /content/reshaped/train/RED WINGED BLACKBIRD
There are 0 folders and 160 image files in /content/reshaped/train/RED HEADED DUCK
There are 0 folders and 166 image files in /content/reshaped/train/BANDED PITA
There are 0 folders and 194 image files in /content/reshaped/train/BANDED BROADBILL
There are 0 folders and 190 image files in /content/reshaped/train/SNOW GOOSE
There are 0 folders and 187 image files in /content/reshaped/train/MCKAYS BUNTING
There are 0 folders and 154 image files in /content/reshaped/train/KAKAPO
There are 0 folders and 204 image files in /content/reshaped/train/JACOBIN PIGEON
There are 0 folders and 150 image files in /content/reshaped/train/FASCATED WREN
There are 0 folders and 156 image files in /content/reshaped/train/CRESTED SHRIKETIT
There are 0 folders and 156 image files in /content/reshaped/train/PEREGRINE FALCON
There are 0 folders and 196 image files in /content/reshaped/train/NORTHERN PARULA
There are 0 folders and 143 image files in /content/reshaped/train/GREY HEADED FISH EAGLE
There are 0 folders and 160 image files in /content/reshaped/train/STRIPPED SWALLOW

Plot pictures from sample species

Plot one picture from 15 randomly selected species

Then I draw a random picture from each class and show their shape and size. A picture's shape shows typically three dimensions. The first two dimensions build a 2D matrix of n rows by m columns. The number of rows represents the image height in pixels, while the number of columns represents the image width in pixels. So each matrix coordinate point represents the intensity value of a single pixel. The third dimension refers to the number of color planes (or channels). There is one plane (2D matrix) for each RGB color, so the value of the third dimension is 3. A picture's size shows its total number of pixels, which results by multiplying the dimensions of the 2D matrices by themselves and by the number of color planes.

```
In [ ]: ### 3) plot one random image from each bird class  
### modules: b_viewClasses  
randomClasses = viewRandomClasses(trainPath, targetClasses[0])  
plt.show()
```

Class: EGYPTIAN GOOSE
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: RED WINGED BLACKBIRD
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: RED HEADED DUCK
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: BANDED PITA
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: BANDED BROADBILL
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: SNOW GOOSE
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: MCKAYS BUNTING
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: KAKAPO
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: JACOBIN PIGEON
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: FASCIATED WREN
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: CRESTED SHRIKETIT
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: PEREGRINE FALCON
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: NORTHERN PARULA
Image shape (rows, columns, channels): (56, 56, 3)
Image size (number of pixels): 9408

Class: GREY HEADED FISH EAGLE

Image shape (rows, columns, channels): (56, 56, 3)

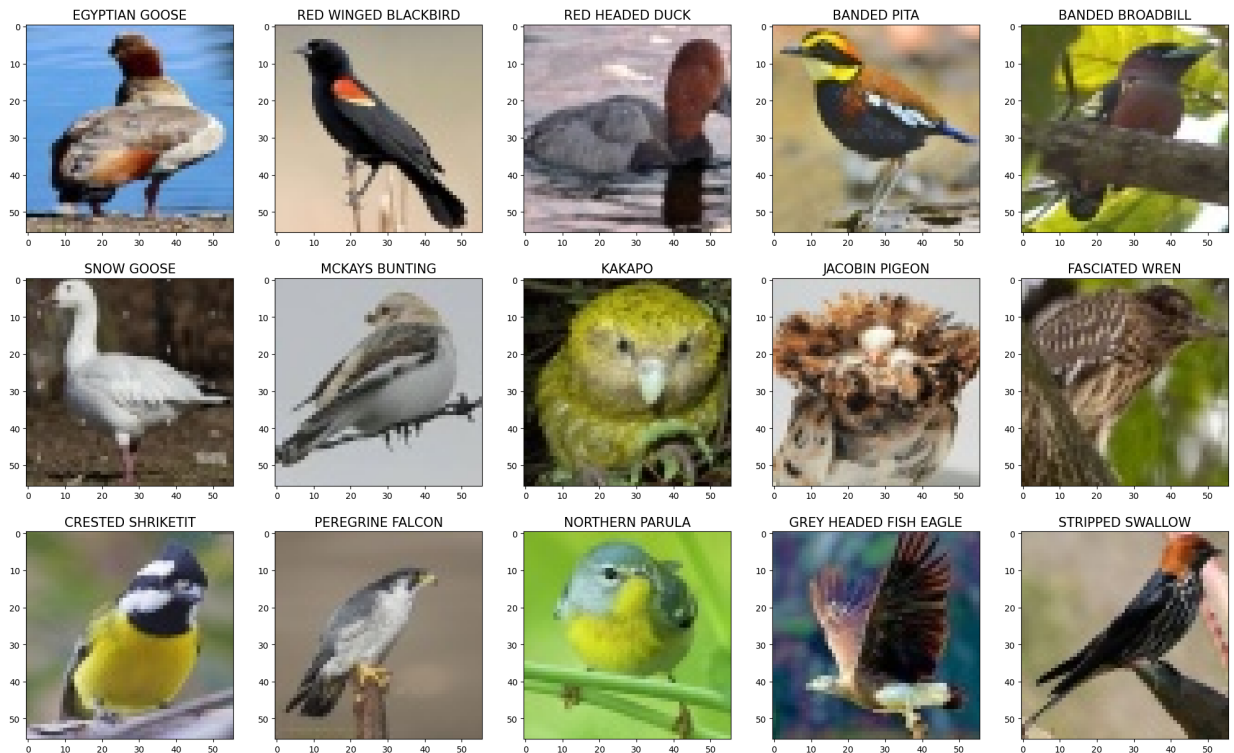
Image size (number of pixels): 9408

Class: STRIPPED SWALLOW

Image shape (rows, columns, channels): (56, 56, 3)

Image size (number of pixels): 9408

Sample of 15 bird classes



Plot 15 pictures from one of the previously selected random species

Finally, we print 15 random images from one of the previously chosen species and show their shape and size as before defined.

```
In [ ]: ### 4) plot 15 random images from one of the classes
### modules: b_viewClasses
randomClass = viewRandomClass(trainPath, targetClasses[0])
plt.show()
print('\nI conclude that, although there is a varying number of images',
      'for each bird class,')
print('in our sample of 15 classes there is a minimum',
      'of', min(targetClasses[2]), 'images, and a maximum of',
      max(targetClasses[2]), 'images.\n')
```

[illegible]

Figure 1 displays a 3x5 grid of 15 images of various bird species, including cockatoos, budgerigars, and lovebirds, used for the bird species classification task. Each image is labeled with a row and column index (e.g., 0, 10, 20, 30, 40, 50) along the axes.

I conclude that, although there is a varying number of images for each bird class, in our sample of 15 classes there is a minimum of 143 images, and a maximum of 204 images.

Data preparation

The following code allows me to load each image of the train and test sets into Python lists. The lists containing the images are then turned into numerical arrays, with 4 dimensions: the first one represents the number of images, while the other three represent the images' shape, which has been rendered homogeneous (56, 56, 3) by removing those not matching the common shape within the execution of the custom *selectData* function.

```
In [5]: ### 5) load the train, test and validation data and labels into memory
#trainData, trainClasses, testData, testClasses, validData, validClasses,
trainData, trainClasses = selectData(trainPath)
testData, testClasses = selectData(testPath)
validData, validClasses = selectData(validPath)
#train, test and validation data lists into numpy arrays
npTrainData = np.array(trainData)
npTestData = np.array(testData)
npValidData = np.array(validData)

print('Our train data array has', npTrainData.ndim, 'dimensions, and a shape of',
      npTrainData.shape, 'for a total number of elements of', npTrainData.size)
print('Our test data array has', npTestData.ndim, 'dimensions, and a shape of',
      npTestData.shape, 'for a total number of elements of', npTestData.size)
print('Our validation data array is equivalent in number to the test data array. It has',
      npValidData.ndim, 'dimensions, and a shape of', npValidData.shape,
      'for a total number of elements of', npValidData.size, '.')
```

Our train data array has 4 dimensions, and a shape of (82724, 56, 56, 3) for a total number of elements of 778267392 .

Our test data array has 4 dimensions, and a shape of (2575, 56, 56, 3) for a total number of elements of 24225600 .

Our validation data array is equivalent in number to the test data array. It has 4 dimensions, and a shape of (2575, 56, 56, 3) for a total number of elements of 24225600 .

The data labels represent the bird species to which the images belong. They are first converted into 2D arrays and finally turned into categorical data: bird classes are substituted by numerical categories.

Bird species should be homogeneous across train and test data, so I check if it's actually the case.


```

In [6]: ### 6) turn the train, test and validation labels into categorical arrays
# train labels
trainClasses = classesToInt(trainClasses)
npTrainClasses = np.array(trainClasses)
npTrainClasses = np.expand_dims(npTrainClasses, axis = 1) # add dimension
npTrainLabels = to_categorical(npTrainClasses)
# test labels
testClasses = classesToInt(testClasses)
npTestClasses = np.array(testClasses)
npTestClasses = np.expand_dims(npTestClasses, axis = 1) # add dimension t
npTestLabels = to_categorical(npTestClasses)
# validation labels
validClasses = classesToInt(validClasses)
npValidClasses = np.array(validClasses)
npValidClasses = np.expand_dims(npValidClasses, axis = 1) # add dimension
npValidLabels = to_categorical(npValidClasses)

trainCount = countLabels(npTrainClasses) # count train labels
testCount = countLabels(npTestClasses) # count test labels
validCount = countLabels(npValidClasses) # count validation labels
if trainCount == testCount and trainCount == validCount:
    print('There are', npTrainClasses.size, 'categories (it means that ea
        'belongs to a category) for a set of', trainCount, 'categories.
else:
    print('ERROR: train labels count, test labels count and validation la

# free up memory
del trainData, testData, validData, trainClasses, testClasses, validClass

```

There are 82724 categories (it means that each image belongs to a category) for a set of 515 categories.

Data normalization

At this point I normalize the data by dividing them for the maximum value they can assume. In this way the data range from 0 to 1 and allow for better speed and prediction results.

```

In [7]: ### 7) normalize the data
npTrainData = np.array(npTrainData / npTrainData.max(), dtype = np.float16)
npTestData = np.array(npTestData / npTestData.max(), dtype = np.float16)
npValidData = np.array(npValidData / npValidData.max(), dtype = np.float16)

```

Call the sequential model

I call the function executing the sequential model from the file d_sequentialModel. The function returns the model history, which is used to train the model and plot the loss function and the model accuracy.

```
In [8]: ### 8) call the sequential model  
myModel, Y_pred = seqModel(npTrainData, npTrainLabels, npTestData, npTest
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
 94765736/94765736 [=====] - 0s 0us/step
 Model: "sequential"

Layer (type)	Output Shape	Param #
resnet50 (Functional)	(None, 2048)	23587712
FC (Dense)	(None, 515)	1055235
Softmax (Activation)	(None, 515)	0

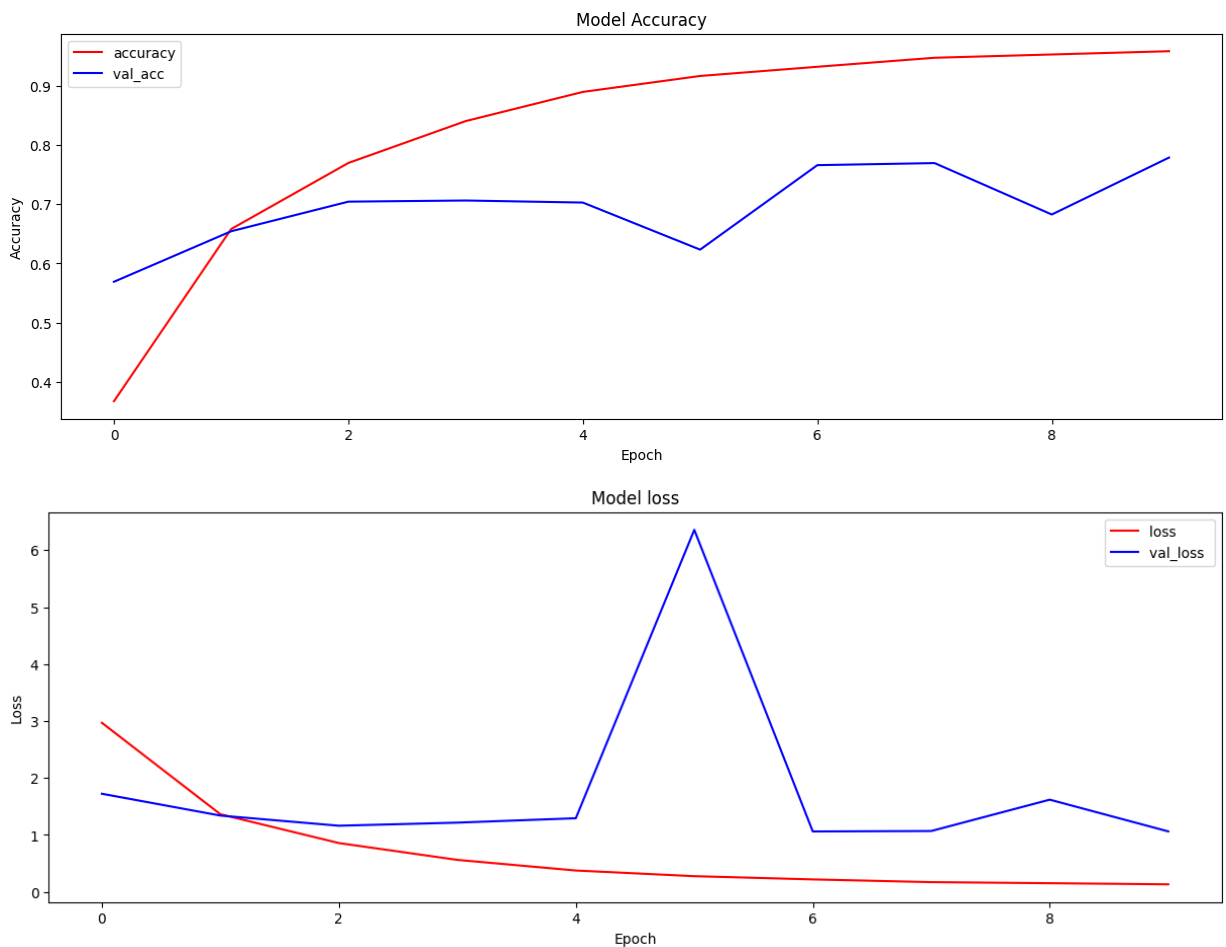
=====
 Total params: 24642947 (94.01 MB)
 Trainable params: 24589827 (93.80 MB)
 Non-trainable params: 53120 (207.50 KB)

Epoch 1/10
 1293/1293 [=====] - 145s 80ms/step - loss: 2.96
 92 - accuracy: 0.3669 - val_loss: 1.7232 - val_accuracy: 0.5689
 Epoch 2/10
 1293/1293 [=====] - 102s 79ms/step - loss: 1.36
 32 - accuracy: 0.6581 - val_loss: 1.3401 - val_accuracy: 0.6544
 Epoch 3/10
 1293/1293 [=====] - 102s 79ms/step - loss: 0.85
 72 - accuracy: 0.7699 - val_loss: 1.1618 - val_accuracy: 0.7045
 Epoch 4/10
 1293/1293 [=====] - 103s 79ms/step - loss: 0.56
 17 - accuracy: 0.8407 - val_loss: 1.2160 - val_accuracy: 0.7064
 Epoch 5/10
 1293/1293 [=====] - 102s 79ms/step - loss: 0.37
 50 - accuracy: 0.8901 - val_loss: 1.2931 - val_accuracy: 0.7029
 Epoch 6/10
 1293/1293 [=====] - 101s 78ms/step - loss: 0.27
 55 - accuracy: 0.9171 - val_loss: 6.3583 - val_accuracy: 0.6233
 Epoch 7/10
 1293/1293 [=====] - 102s 79ms/step - loss: 0.21
 87 - accuracy: 0.9326 - val_loss: 1.0609 - val_accuracy: 0.7662
 Epoch 8/10
 1293/1293 [=====] - 102s 79ms/step - loss: 0.17
 09 - accuracy: 0.9479 - val_loss: 1.0688 - val_accuracy: 0.7697
 Epoch 9/10
 1293/1293 [=====] - 101s 78ms/step - loss: 0.15
 22 - accuracy: 0.9534 - val_loss: 1.6195 - val_accuracy: 0.6827
 Epoch 10/10
 1293/1293 [=====] - 102s 79ms/step - loss: 0.13
 25 - accuracy: 0.9589 - val_loss: 1.0618 - val_accuracy: 0.7790

Test loss: 0.7936546802520752
 Test accuracy: 0.8229126334190369
 81/81 [=====] - 2s 13ms/step

I finally plot the loss function for the model and the model accuracy.

```
In [9]: ### 9) plot model accuracy and loss function  
# accuracy  
plt.figure(figsize=(15, 5))  
plt.plot(myModel.history['accuracy'], 'r', label='accuracy')  
plt.plot(myModel.history['val_accuracy'], 'b', label='val_acc ')  
plt.title('Model Accuracy')  
plt.ylabel('Accuracy')  
plt.xlabel('Epoch')  
plt.legend()  
  
# loss  
plt.figure(figsize=(15, 5))  
plt.plot(myModel.history['loss'], 'r', label='loss ')  
plt.plot(myModel.history['val_loss'], 'b', label='val_loss ')  
plt.title('Model loss')  
plt.ylabel('Loss')  
plt.xlabel('Epoch')  
plt.legend()  
  
plt.show()
```



Plot confusion matrix and classification report

I finally plot the confusion matrix and the classification report for a small sample of bird species.

```
In [10]: ### 10) Confution Matrix and Classification Report
y_pred = np.argmax(Y_pred, axis=1)
cm = confusion_matrix(npTestClasses, y_pred)
thresh = cm.max() / 2.
tick_marks = np.arange(len(classes))
target_names = classes
print('\nConfusion Matrix (small sample)\n')
print(cm[0 : 15, 0 : 15])
plotCM(cm[0 : 15, 0 : 15], target_names[0 : 15])
print('\n\nClassification Report (small sample)\n')
class_report = classification_report(npTestClasses, y_pred, target_names
print(class_report[0 : 1142], (' ' * 18) + ' [...] \n') # print a few lines
```

Confusion Matrix (small sample)

```

[[5 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 5 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 5 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 4 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 4 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 3 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 5 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 4 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 5 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 3 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 5 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 5 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 5 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 5 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 0 5]]

```

Classification Report (small sample)

	precision	recall	f1-score	support
BLACK-CAPPED CHICKADEE	1.00	1.00	1.00	5
BLONDE CRESTED WOODPECKER	1.00	1.00	1.00	5
WHITE BROWED CRAKE	0.62	1.00	0.77	5
AMERICAN GOLDFINCH	0.80	0.80	0.80	5
LIMPKIN	0.80	0.80	0.80	5
APAPANE	0.75	0.60	0.67	5
BLACKBURNIAM WARBLER	1.00	1.00	1.00	5
NORTHERN FLICKER	0.80	0.80	0.80	5
PAINTED BUNTING	1.00	1.00	1.00	5
AUSTRALASIAN FIGBIRD	0.60	0.60	0.60	5
ANDEAN SISKIN	0.56	1.00	0.71	5
EMPEROR PENGUIN	0.56	1.00	0.71	5
TRICOLORED BLACKBIRD	1.00	1.00	1.00	5
RUBY THROATED HUMMINGBIRD	1.00	1.00	1.00	5
PHILIPPINE EAGLE	1.00	1.00	1.00	5
[...]				

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

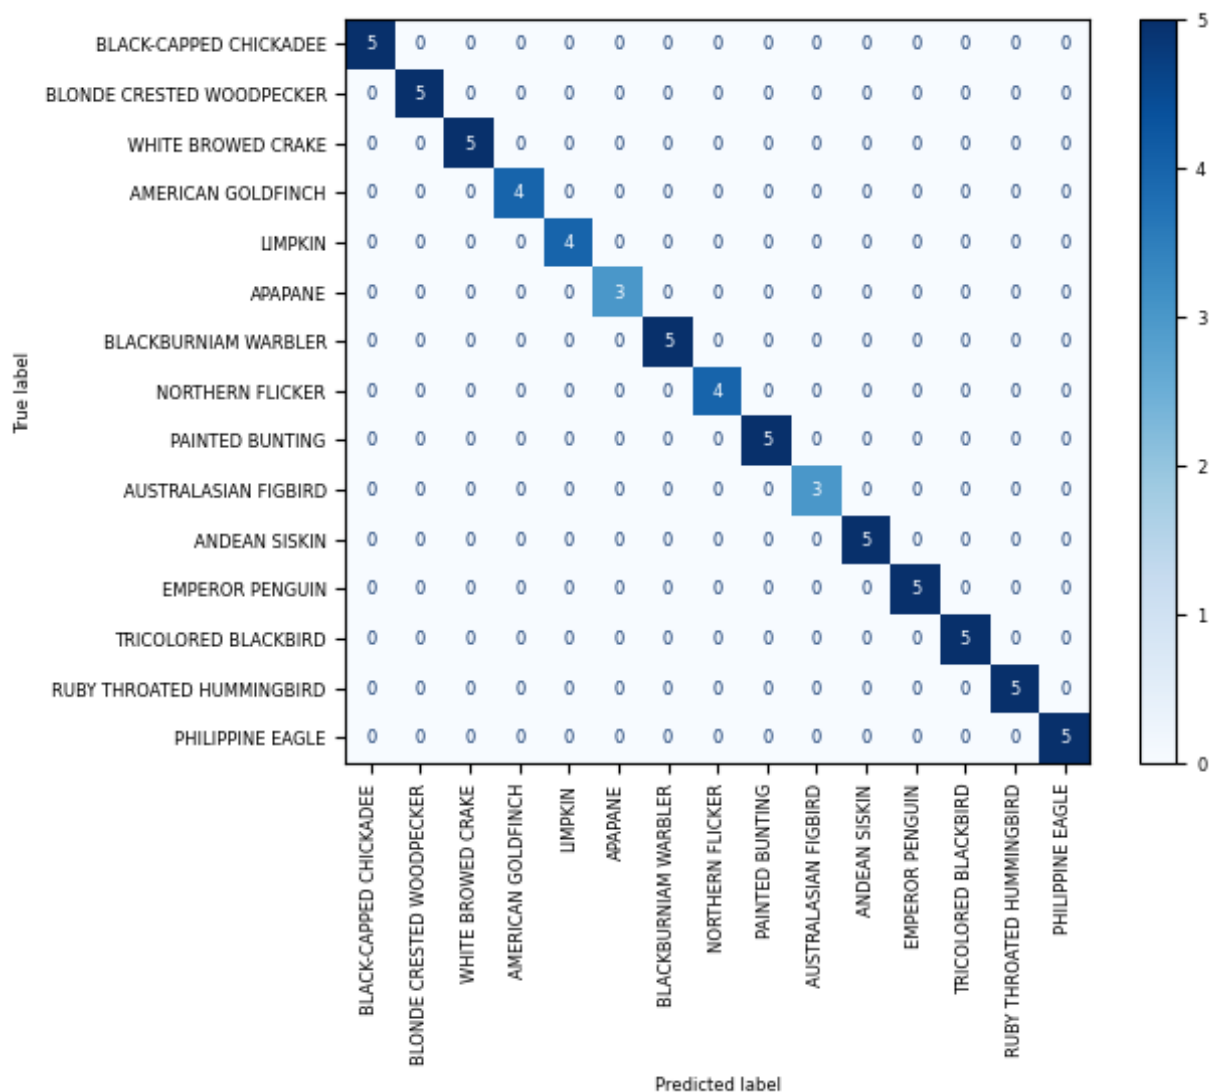
```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

<Figure size 600x600 with 0 Axes>



In []: